

INSTRUCTOR GUIDE

Organizational Design for Product Teams

Course 04 — Facilitation Guide for Graduate Assistants and Faculty

Patrick T. Hoffman · @arcktip · patrickthoffman.com

HOW TO USE THIS GUIDE

This guide supports graduate teaching assistants and faculty facilitating Course 04: Organizational Design for Product Teams. Each session entry provides a Session Overview, Theoretical Frame, Discussion Arc with timing, numbered Discussion Questions for all students (with facilitation notes), Graduate Questions (G1–G2, teal-accented), Common Misconceptions, Cross-session Connections, and Facilitator Notes.

The course is designed for a dual-listed seminar (upper-division undergraduate + master's) drawing students from urban planning, public policy, MBA, and engineering programs. The graduated question structure makes the two levels separable if needed.

Patrick's practitioner voice in this course

Course 04 draws directly on Patrick Hoffman's experience designing and redesigning product organizations across every stage of growth — from 0-to-10 PMs at Social Bicycles and JUMP, through product leadership at Google and Meta, to the organizational challenges of Udacity and independent consulting. The Rule of Seven framework (@arcktip) is the course's primary original contribution. Week 11 (executive guest session) is ideally anchored by Patrick or a senior product leader from his network. His account of the JUMP/Uber integration (Week 12) is the primary practitioner case for mergers and reorgs.

Final proposal evaluation criteria — share at Week 1

Students build toward these criteria from the first session. The final org design proposal must address all six:

- Diagnostic rigor — is the current-state pathology traced to a specific structural cause using course frameworks?
- Rule of Seven compliance — does the proposed structure satisfy the span-of-control principles?
- Mission clarity — does every team in the proposed structure have a mission that passes the five-sentence test?
- Conway's Law alignment — does the proposed structure produce the products the company needs?
- Implementation realism — is there a stakeholder map, a coalition strategy, and an early-win sequence?
- Trade-off honesty — are the costs of the proposed redesign explicitly named, not buried or

minimized?

Week 1 · Organizations as Products

FORMAT

Seminar

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

The opening session installs the course's central reframe: an organization is a product, and like any product, it can be well or poorly designed. A poorly designed org ships poorly designed products — or doesn't ship at all. A well-designed org ships on time, makes good decisions at the right level, and retains the people who make it work. Most students arrive having experienced organizational dysfunction without a vocabulary for diagnosing it. This session gives them that vocabulary and immediately applies it to cases they already know.

THEORETICAL FRAME

Conway's Law is the session's conceptual anchor: 'Any organization that designs a system will produce a design whose structure is a copy of the organization's communication structure.' Published in 1968, it remains the most predictive single insight in org design. Galbraith's Star Model (strategy, structure, processes, rewards, people) provides the full diagnostic framework. Cagan and Jones's 'Empowered' grounds both in contemporary product management practice.

SUGGESTED DISCUSSION ARC

0–10 min

Open with a product failure the class knows: the Google Stadia shutdown, the Meta Portal discontinuation, or the WeWork IPO collapse. 'Don't tell me about the product. Tell me about the organization that produced it.' Surface: every product failure is also an organizational failure — the product reflected what the org was capable of building and deciding.

10–25 min

Lecture: Conway's Law in practice. Show two org charts side by side — a siloed, functional org and a cross-functional, product-team org — then show the products each typically produces. The functional org produces products with visible seams at the points where functions hand off to each other. The cross-functional org produces products that feel integrated. This is Conway's Law.

25–50 min

Seminar discussion: Questions 1–3. Question 2 (what can and cannot be fixed by changing the org) is the most important conceptual question of the course — return to it in every subsequent session.

50–65 min

Graduate section: Galbraith, Lawrence/Lorsch, Chandler, Mintzberg. Questions G1 and G2. The Chandler question (structure follows strategy, or strategy follows structure?) is the session's deepest theoretical debate.

65–75 min

Close: each student names one organizational dysfunction they have personally experienced. Collect these. 'This is the diagnostic data set for the course — we will return to your examples all semester.' Keep the list visible.

DISCUSSION QUESTIONS — ALL STUDENTS

- 1** Conway's Law says that a system's design will mirror the organization that built it. Find a product you use regularly and work backward: what does the product's design reveal about the organizational structure of the team that built it?

Best examples: the seams in any Microsoft product reveal functional silos. The seamlessness of an early iPhone reflects a highly integrated team. Any product with an inconsistent UX across features has a fragmented PM structure. Students with product experience can often identify their own organization's products this way.

- 2** What are the things that organizational design can fix, and what are the things it cannot? Draw the line.

Org design can fix: coordination costs, decision rights ambiguity, span of control overload, mission confusion, platform bottlenecks. Org design cannot fix: bad strategy, bad people, insufficient resources, or market conditions. The distinction matters enormously — a reorg that is actually a strategy pivot in disguise will fail.

- 3** Cagan and Jones argue that the difference between an 'output team' (a team that ships features) and an 'outcome team' (a team that owns business results) is fundamentally an organizational design question, not a PM mindset question. Do you agree? What would a company need to change structurally to move from output to outcome teams?

The key structural changes: outcome teams need a measurable outcome they own, decision-making authority over how to achieve it, and a manager who evaluates them on outcomes not output. None of those changes is a mindset shift — they are org design decisions made by senior leadership.

GRADUATE / MASTER'S QUESTIONS

- G1** Chandler argued 'structure follows strategy.' Lawrence and Lorsch argued that structure depends on environmental uncertainty — more uncertain environments require more differentiation and more integration simultaneously. How do these two frameworks interact? Which better predicts the org design choices of a fast-growing tech company in 2026?

- G2** Mintzberg describes five configurations of organizations: the simple structure, machine bureaucracy, professional bureaucracy, divisionalized form, and adhocracy. Which configuration best describes a 50-person product startup? A 5,000-person tech company? Does the transition between configurations explain 'scaling pains'?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students often read Conway's Law as prescriptive (org should mirror product architecture) rather than descriptive (org does mirror product architecture, whether you intend it to or not). The prescription comes second: if you want to change the product structure, you must change the org structure first.

The 'output vs. outcome teams' framing is sometimes heard as a critique of engineers and PMs who lack ambition. The argument is structural: a team cannot own outcomes it doesn't have the authority to influence. Blame the org design, not the team.

Executive MBA students sometimes arrive with a bias toward the 'simple structure' startup model as inherently superior. The course's argument is more nuanced: the right org design depends on the strategy, the environment, and the stage of the company. There is no universally correct structure.

CONNECTIONS TO OTHER SESSIONS

Conway's Law is the course's most cited single concept. Return to it explicitly every time a new org structure is introduced: 'What products does this structure produce, and does Conway's Law predict those products?' The Week 6 (platform vs. product teams) session is the deepest application of Conway's Law in the course.

FACILITATOR NOTES

The product-failure-as-org-failure opening works best with a failure that is recent, visible, and that students feel they understand. Stadia is good for tech-industry students. WeWork is good for MBA students. The key move: every student in the room has a failure story from their own work experience. Invite one or two to share it before revealing the Stadia case — the personal examples are more instructive than the famous ones.

Week 2 · The Rule of Seven

FORMAT

Lecture + workshop

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

The Rule of Seven is the course's primary original framework and the analytical tool students will apply most frequently in subsequent sessions and in their own careers. The rule: no manager should have more than seven direct reports who are themselves managers; the optimal span for knowledge work teams is five to seven. This session teaches the rule's foundations (Dunbar's cognitive limits, Miller's magical number, team research from Hackman) and immediately applies it to org charts students bring from their own workplaces.

THEORETICAL FRAME

The Rule of Seven is not a rigid constraint but a diagnostic tool: when a manager has twelve direct reports, something is being sacrificed — either the quality of individual attention, the depth of strategic engagement, or the manager's own capacity for the work that only they can do. When a manager has two, something is also wrong — excessive hierarchy, insufficient trust, or a structure that hasn't caught up to a team that has shrunk.

SUGGESTED DISCUSSION ARC

0–10 min

Open with the math: if a company has 1,000 employees and a manager can effectively manage 7 direct reports, how many layers of management does the company need? (Log base 7 of 1,000 $\approx$ 3.6 layers.) Now what if the span is 3? (Log base 3 of 1,000 $\approx$ 6.3 layers.) 'The difference between a span of 3 and a span of 7 is more than 2 layers of management in a 1,000-person company. Every one of those layers is a decision delay and a communication loss.'

10–25 min

Lecture: the three foundations of the Rule of Seven. (1) Dunbar's number: humans can maintain stable relationships with approximately 150 people; closer relationships require more cognitive investment and top out around 5-7 for genuine mentorship. (2) Miller's law: working memory can hold roughly 7 plus-or-minus 2 chunks — the same constraint applies to managing the complexity of a team. (3) Hackman's team conditions: real teams require a compelling direction, an enabling structure, a supportive context, and expert coaching — each requires invested managerial attention that is finite.

25–55 min

Workshop: students draw their own org chart (real or from memory) and audit it against the Rule of Seven. For each manager: count their direct reports, assess whether they are managing managers or ICs (the Rule applies differently to each), and identify where span of control is too wide or too narrow. Questions 1–3 weave through the workshop.

55–70 min

Graduate section: Hackman's Leading Teams, Dunbar, Miller, Ancona/Caldwell, Belbin. Questions G1 and G2.

70–75 min

Close: each student identifies the single manager in their audit whose span of control is most egregiously wrong — too wide or too narrow. 'That manager's span is causing a specific organizational problem. Name it.'

DISCUSSION QUESTIONS — ALL STUDENTS

- 1 **A VP of Product at a 200-person company has 12 direct reports, all senior product managers. The VP complains that they can never finish strategic work because they spend all day in one-on-ones and check-ins. Apply the Rule of Seven: what is the diagnosis, and what is the structural fix?**
Diagnosis: the VP's span of control is too wide — 12 is nearly double the upper bound. The VP is doing the work of two or three managers, sacrificing strategic work for managerial overhead. Fix: create a layer of Group Product Managers (GPMs) between the VP and the PMs, each with a span of 4–5, freeing the VP to manage 3–4 GPMs and do actual product strategy.
- 2 **The Rule of Seven applies differently when a manager's direct reports are themselves managers vs. when they are individual contributors. Why? What should the span of control be for a manager of managers, and how does it differ from the span for a manager of ICs?**
Manager of managers: lower span (4–6) because each direct report requires significant coordination, context, and coaching investment. Their work is organizational and complex. Manager of ICs: can support slightly higher span (6–8) for routine work, lower (4–6) for highly creative or highly uncertain work. The span should match the cognitive and coordination demand of the managed work.
- 3 **A company has a policy that all engineering managers must have at least 8 direct reports for the management role to be 'justified.' What does that policy reveal about how the company thinks about management? What is it likely producing?**
The policy treats management as a cost to be minimized (managers should be 'used efficiently'). It produces: over-stretched managers who do insufficient coaching, PMs and engineers who don't receive enough direction, and either invisible performance problems (no one has time to catch them) or burnout (everyone is overloaded). The policy implicitly assumes that management is overhead rather than leverage.

GRADUATE / MASTER'S QUESTIONS

- G1 Hackman identifies six conditions for team effectiveness: being a real team, having a compelling direction, an enabling structure, a supportive context, expert coaching, and adequate information. How do these conditions map onto span of control? At what span does each condition become difficult to maintain?
- G2 Miller's 'magical number seven' applies to working memory for information chunks. Dunbar's number applies to stable social relationships. Both converge on a similar limit. Is this convergence evidence of a deep constraint, or is it coincidence? What does either answer imply for how we should think about org design?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

The Rule of Seven is frequently misapplied as 'everyone must have exactly seven reports.' The rule is a range and a principle: spans below 4 and above 8 are almost always suboptimal, but the specific number depends on the nature of the work, the seniority of the team, and the manager's own responsibilities beyond management.

Students sometimes argue that great managers can handle 15+ direct reports. This is true — briefly, and at great personal cost. The question is not whether a great manager can survive a too-wide span; it's whether the organization can sustain that dependency and whether the manager can do any work beyond managing.

Graduate students may treat the cognitive science basis of the rule (Dunbar, Miller) as foundational proof. It is suggestive, not definitive. The rule's primary support is empirical: organizations that consistently violate it exhibit predictable pathologies. The cognitive science provides a plausible mechanism, not a mathematical proof.

CONNECTIONS TO OTHER SESSIONS

The Rule of Seven is applied in every subsequent session. In Week 4 (startup org design), it determines when a PM team needs a first manager. In Week 5 (scaling pains), violations of the rule are the primary diagnostic. In Week 10 (enterprise org design), it determines the depth of the management hierarchy. In Week 15 (final presentations), every org design proposal is audited against it.

FACILITATOR NOTES

The math opening (layers of management at different spans) is highly effective with MBA audiences — they respond to quantitative framing. Have students do the calculation themselves rather than showing it. The org chart audit requires that students bring an actual org chart, or can reconstruct one from memory. Assign this explicitly before class: 'Come prepared to draw your team's org chart on paper.'

Week 3 · Mission-Driven Teams

FORMAT

Seminar + writing workshop

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

The team mission is the most underestimated tool in org design. A mission that is written well creates alignment without mandating methods, focuses the team on outcomes without specifying solutions, and remains stable enough to guide decisions through a 12-month product cycle. A mission written poorly — too broad, too tactical, or too vague — creates exactly the ambiguity it was supposed to resolve. This session teaches students to write missions, critiques missions from real product teams (with identifying information removed), and establishes the five-sentence mission test from the @arcktip framework.

THEORETICAL FRAME

The tension between autonomy and coordination is the session's intellectual engine. Full autonomy (every team does what it thinks is right) produces incoherence at the portfolio level. Full coordination (every team is told what to do) produces slow, unempowered execution. The well-written mission is the mechanism that resolves this tension: it tells the team what to achieve and why, without telling them how. OKRs are the operating system that sits on top of the mission, connecting it to the quarterly rhythm.

SUGGESTED DISCUSSION ARC

0–10 min

Show three product team missions. Mission A: 'Grow revenue for the payments product.' Mission B: 'Build a seamless checkout experience that earns the trust of merchants and buyers globally.' Mission C: 'Help people pay.' Ask: 'Which team would you rather be on? Which team would you rather manage? Which team has the clearest decision-making compass?' Surface: Mission A is an output measure, not a mission. Mission C is so broad it provides no guidance. Mission B is constraining enough to focus decisions and flexible enough to allow methods.

10–20 min

Lecture: the five-sentence mission test (@arcktip). (1) A good mission specifies who the team serves. (2) It specifies what the team enables or prevents. (3) It is testable: a year from now, you can evaluate whether the mission was accomplished. (4) It is exclusive: it rules out work as clearly as it rules it in. (5) It survives at least one 'why' — it connects to a larger organizational purpose without requiring elaboration.

20–50 min

Writing workshop: students draft or critique a mission for a real team. Questions 1–3 weave through the workshop. The critique activity (applying the five-sentence test to a mission someone else wrote) is more productive than writing from scratch.

50–65 min

Graduate section: Simons, Collins/Porras, Edmondson (Teaming), Grant. Questions G1 and G2.

65–75 min

OKR connection: brief lecture on how the mission anchors the OKR hierarchy. The mission is not an OKR — it is the constant that OKRs are measured against. 'Your mission should be true whether you hit your Q3 OKRs or not.'

DISCUSSION QUESTIONS — ALL STUDENTS

- 1 **The difference between a mission and a goal is that a mission is directional and durable, while a goal is specific and time-bound. What happens when a product team uses a goal as its mission? What behavior does that produce?**
When a goal becomes a mission: the team optimizes for hitting the goal rather than for the underlying purpose. Classic example: 'Grow DAU to 10M' as a team mission produces engagement hacks that inflate the metric without serving users. The mission should specify why DAU matters, not the DAU target itself.
- 2 **Apply the five-sentence mission test to the mission of a team you know — your current employer, a past employer, or a case from the course. Which of the five criteria does the mission fail? What would you change?**
Invite 2–3 students to share their missions. Apply the test publicly. The most common failure: missions that specify outputs (build X) rather than outcomes (enable Y for Z users). The second most common: missions so broad they include everything ('improve the customer experience' could mean anything).
- 3 **Two adjacent teams in the same product org have missions that overlap — both claim ownership of the customer onboarding experience. What organizational failure does that overlap reveal? How do you resolve it without a reorg?**
The overlap reveals a failure of portfolio design: the org did not define team boundaries precisely enough when the teams were formed. Resolution without a reorg: explicit decision-rights assignment (one team owns the experience; the other owns the infrastructure) and a coordination mechanism (regular joint review) until the next structural review.

GRADUATE / MASTER'S QUESTIONS

- G1 **Simons's levers of control framework distinguishes between belief systems (mission, values), boundary systems (rules and limits), diagnostic systems (metrics and feedback), and interactive systems (strategy conversations). How does a team mission relate to each of these four levers? Which lever does a poorly written mission fail to activate?**
- G2 **Edmondson's 'Teaming' describes organizations where team membership is fluid — people form and dissolve teams rapidly based on project needs. In a teaming organization, what replaces the team mission as the mechanism of coordination? Is a teaming org compatible with the Rule of Seven?**

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

The mission writing exercise often surfaces a common misconception: that a mission should be inspirational above all else. An inspiring mission that fails the five-sentence test creates emotional commitment to a vague direction — which is often worse than no direction, because it generates confident action without guidance.

OKRs are frequently confused with missions. The mission is the 'why' that stays stable for 2–5 years; OKRs are the 'what and how' for the next 90 days. A team whose mission changes every quarter has no mission — only a series of goals.

Graduate students may push back that Simons's control systems are more relevant to large enterprises than to product startups. The response: startups with weak belief systems and no boundary systems produce the worst governance failures (Theranos, WeWork). The frameworks apply at any scale; the implementation differs.

CONNECTIONS TO OTHER SESSIONS

The mission-driven team framework is applied in the midterm case (Week 8) and the final presentation (Week 15). In Week 9 (cross-functional collaboration), mission overlap is the primary source of conflict. In Week 14 (leading the redesign), rewriting team missions is the most politically charged step in an org redesign.

FACILITATOR NOTES

The three-missions opening works best if you can show real missions from real companies — scrubbed of identifying information if necessary. If you cannot source real missions, the three constructed examples in the arc work well. The writing workshop should produce a real output that students keep: their own team's mission, improved. That output has practical value beyond the course.

Week 4 · Startup Org Design

FORMAT

Case seminar

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

Startup org design is the most consequential and least discussed dimension of early-stage company building. Most startup advice focuses on product-market fit, fundraising, and hiring — and says almost nothing about when to add a first manager, how to structure the engineering-PM-design triad at 20 people, or how to avoid the premature bureaucratization that kills velocity. This session uses the 0-to-10 PM arc as a spine: each new PM hire changes the organizational problem the founder or CPO faces.

THEORETICAL FRAME

Horowitz's account of the CEO's organizational role is the most honest practitioner text in this space — he writes about the specific decisions (when to promote from within vs. hire from outside, when to add a layer of management, how to handle a manager who is loyal but no longer scaling) that most org design books treat as self-evident. Hoffman's @arcktip account of the SoBi/JUMP arc provides a mobility-product-specific case that students will recognize from earlier courses.

SUGGESTED DISCUSSION ARC

0–10 min

Open with a progression: '1 PM, no structure. 3 PMs, still no structure. 6 PMs, starting to feel it. 10 PMs, you have an organizational problem.' Ask: 'What breaks at each step? What is the first thing that breaks at 6 PMs that wasn't broken at 3?' Surface: at 3, informal coordination still works. At 6, you need explicit decision rights, a shared planning process, and someone managing the PMs — whether or not the org chart reflects that yet.

10–25 min

Lecture: the 0-to-10 PM arc. (1) 0–2 PMs: the founder-PM model — the product leader is still the PM. Key decision: when to make the first PM hire, and what type of PM to hire first. (2) 3–5 PMs: the peer team — coordination is still informal. Key decision: who sets priorities when PMs disagree? (3) 6–8 PMs: the coordination crisis — informal mechanisms break. Key decision: first manager or flat VP? (4) 9–12 PMs: the first management layer — now you are managing managers. Key decision: what is the manager-to-PM ratio, and who is the right first manager?

25–55 min

Case seminar: students analyze a startup org design case — ideally a company they know well. Questions 1–3.

55–70 min

Graduate section: Blank, Eisenmann, Feld/Mendelson, Catmull. Questions G1 and G2.

70–75 min

Close: 'What is the earliest org design decision you have ever seen a startup get wrong? What did it cost?' Rapid sharing — 2 minutes per student, no more than 5 students.

DISCUSSION QUESTIONS — ALL STUDENTS

- 1** A startup has 5 PMs, all reporting to the co-founder who is also the CPO. The co-founder is spending 60% of their time in one-on-ones and product reviews. What is the organizational diagnosis, and what is the specific decision the co-founder needs to make?

Diagnosis: the co-founder's span of control is wide relative to the organizational work they are being asked to do. At 5 PMs, this is borderline — some co-founders can manage 5–6 ICs effectively. But if the co-founder's strategic work is suffering, the diagnosis is unambiguous: they need either a first manager layer or to hire a Head of Product to hold the direct reports while they step up to CEO-level strategic work.

- 2** Horowitz distinguishes between 'the peacetime CEO' (who can optimize and develop talent in a stable environment) and 'the wartime CEO' (who must act decisively in a crisis). How does that distinction apply to org design decisions at a startup? When does a startup need a 'wartime' org structure, and what does that look like?

Wartime org structure: flatter (fewer layers), more centralized decision-making, faster but potentially worse decisions with less consultation. Peacetime: more layers, more delegation, slower decisions with more buy-in. Most startups oscillate between the two — the org design must accommodate both modes, or the company will restructure every six months.

- 3** A startup has its first PM hire start on Monday. The CEO asks you: 'What are the three most important things we should decide about this person's role before they arrive?' What do you say?

The three decisions: (1) What does this PM own — which product surface, which user problem, which team? (2) Who has final decision authority when the PM and an engineer disagree? (3) What does 'success in 90 days' look like — and is it defined in outcome terms or output terms?

GRADUATE / MASTER'S QUESTIONS

- G1** Eisenmann's research on why startups fail identifies 'premature scaling' as one of the most common causes — building out org structure, processes, and headcount before achieving product-market fit. How does premature structure manifest in a product organization? What are the early warning signs?
- G2** Catmull's account of Pixar describes a creative organization that maintained a flat, director-driven structure even as it scaled to hundreds of employees. What organizational design principles made that possible? Are those principles generalizable to a product org outside the creative industries?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students often assume that the correct solution to 'coordination problems at 6 PMs' is 'add a manager.' Sometimes the correct solution is 'simplify the portfolio' — fewer PMs each owning clearer domains is

often more effective than adding a management layer to coordinate complex overlapping domains.

The 'wartime vs. peacetime CEO' distinction is sometimes read as: wartime = bad, peacetime = good. Horowitz's point is that both modes are necessary at different moments. A CEO who can only operate in one mode will fail in the other.

Graduate students may push back on the 0-to-10 PM arc as too linear. Real startup org design is messier — companies shrink as well as grow, PMs get promoted and fail, org structures get imposed by investors. Acknowledge the messiness while maintaining that the arc describes the typical trajectory that most startups experience.

CONNECTIONS TO OTHER SESSIONS

The startup org design arc connects directly to Week 5 (scaling pains — what the arc looks like when things go wrong) and Week 8 (midterm case — students often choose a startup or growth-stage company for their midterm). The first-manager decision connects to Week 13 (mentorship as leverage — the first manager's primary job is developing the PMs, not just directing them).

FACILITATOR NOTES

This session works especially well when students have startup experience and can share specific org design decisions they made or witnessed. Executive MBA students often have the richest material. Give permission for confidential sharing: 'You can describe the decision without naming the company.' The most instructive cases are always the failures, not the successes.

Week 5 · Scaling Pains

FORMAT

Case seminar + diagnostic

TIMING

90 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

Scaling pains are the most common and most predictable organizational failure mode in the tech industry. They occur when a company's organizational structure fails to keep pace with its growth — headcount increases, but coordination mechanisms, decision rights, and management quality do not scale with it. This session uses the Spotify squad model as the primary case (a model that was widely adopted by companies that did not understand its preconditions) and draws on McChrystal's JSOC transformation as the deepest account of organizational redesign under pressure.

THEORETICAL FRAME

The Spotify model is the course's most important case study precisely because it was so widely misapplied. Spotify's squads, tribes, chapters, and guilds emerged from a specific engineering culture, at a specific scale, with specific leadership characteristics. Companies that copied the model without the culture got the vocabulary without the function. McChrystal's transformation of JSOC illustrates the same principles — shared consciousness and empowered execution — in a radically different context, which tests whether they are genuinely general.

SUGGESTED DISCUSSION ARC

0–10 min

Ask: 'Has anyone in this room worked at a company that adopted the Spotify model?' If yes: 'Did it work?' The answers are almost always illuminating. If no: 'Have you heard it recommended? What was the appeal?' Surface: the Spotify model is appealing because it sounds like a solution to the coordination problem that everyone experiences at scale. It is not a solution — it is an organizational design philosophy that only works under specific conditions.

10–25 min

Brief lecture: Greiner's five phases of organizational growth and their characteristic crises. (1) Growth through creativity → crisis of leadership. (2) Growth through direction → crisis of autonomy. (3) Growth through delegation → crisis of control. (4) Growth through coordination → crisis of red tape. (5) Growth through collaboration → next crisis unknown. 'Where is your company on this arc? What crisis are you in?'

25–60 min

Case analysis: the Spotify model. What were the conditions that made squads/tribes/chapters work at Spotify? Which of those conditions were missing at the companies that tried to copy it? Questions 1–3.

60–80 min

Graduate section: McChrystal (deep), Greiner, Pfeffer, Sutton. Questions G1 and G2. The McChrystal question (shared consciousness as a precondition for empowered execution) is the session's deepest analytical moment.

80–90 min

Diagnostic: students apply Greiner's model to their own organization or case. 'Which

phase is the organization in? What crisis is it heading toward? What org design intervention would extend the current phase?

DISCUSSION QUESTIONS — ALL STUDENTS

- 1** The Spotify model has been adopted by hundreds of companies. Most of them describe the adoption as unsuccessful. What were the specific conditions at Spotify in 2012 that made the squad model work? Which of those conditions are present at a company you know? Which are absent?

Conditions at Spotify: a relatively flat, highly skilled engineering culture with strong individual autonomy norms; a leadership team that genuinely delegated decision-making; a product architecture (microservices) that enabled team autonomy by reducing inter-team dependencies; and a culture of psychological safety that made the informal coordination work. Most companies copy the vocabulary and the org chart, not the conditions.
- 2** McChrystal describes the JSOC transformation as moving from an 'efficiency' model (optimized chains of command, specialized functions) to an 'adaptability' model (shared consciousness across the network, empowered execution at the team level). What was the specific organizational mechanism that created shared consciousness in JSOC? Can that mechanism be replicated in a product company?

The mechanism: the daily Operations and Intelligence Brief (O&I) — a 2-hour all-hands call, seven days a week, that gave every unit in the network visibility into what every other unit knew and was doing. In a product company: the equivalent is a weekly all-hands with real operational transparency — not a polished update, but genuine information sharing across teams about what is working and what is not.
- 3** A product organization has 40 PMs, 150 engineers, and 20 designers. The CEO reports that every major product decision takes at least three months to make. Apply Greiner's framework: which growth phase is this company in? What is the organizational intervention?

Likely in Phase 3 (growth through delegation) entering the crisis of control, or Phase 4 (growth through coordination) with excessive red tape. Three-month decision cycles at this scale suggest: too many approval layers, unclear decision rights, or a management culture that escalates decisions rather than making them. Interventions: RAPID framework for decision rights, clearer team mission boundaries, and an explicit decision log so teams can learn from each other.

GRADUATE / MASTER'S QUESTIONS

- G1** McChrystal argues that 'empowered execution without shared consciousness is dangerous.' Apply that principle to a product organization where teams are given autonomy before information flows are established. What specific failure modes emerge? What is the minimum viable shared consciousness for a product organization of 200 people?
- G2** Pfeffer's analysis of organizational power suggests that formal org chart changes rarely produce the power redistribution they intend, because informal power structures are more durable than formal ones. What does that imply for org redesigns that attempt to empower lower-level teams? How do

you move informal power?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

The Spotify model is not wrong — it worked at Spotify, and elements of it work in other contexts. The mistake is treating it as a general solution rather than a specific instance of broader organizational principles. The course's goal is to help students derive those principles, not copy the model.

Greiner's phases are sometimes read as inevitable and irreversible. In practice, companies can extend phases through deliberate org design interventions, and can temporarily move backward during crises or contractions. The model is a map, not a destiny.

Students sometimes confuse 'scaling pains' with 'growth problems' — as if pain were a sign of success. Some scaling pains are necessary transitions; others are avoidable failures of foresight. The diagnostic skill is knowing which is which.

CONNECTIONS TO OTHER SESSIONS

This is the course's heaviest case session — the diagnostic skills developed here are directly applied in the midterm (Week 8) and the final presentation (Week 15). The McChrystal connection (shared consciousness) recurs in Week 9 (cross-functional collaboration) and Week 14 (leading the redesign). The Greiner framework is the scaffolding for the Week 10 enterprise org design session.

FACILITATOR NOTES

If any student has personal experience with a Spotify model adoption, invite them to share it before the case discussion. Even one firsthand account is more instructive than the case study. If none are available, the Kniberg/Ivarsson presentation slides (freely available online) can be shown and analyzed in real time — they are the primary source and are unusually honest about what Spotify does not claim the model can do.

Week 6 · Platform vs. Product Teams

FORMAT

Seminar + design exercise

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

The platform team is the most misunderstood structure in product organization design. When done well, a platform team enables product teams to move faster by providing shared infrastructure they don't have to build themselves. When done poorly — which is most of the time — the platform team becomes a bottleneck that slows every team down, demands roadmap prioritization from all of them, and creates the illusion of efficiency while producing the reality of dependency. This session teaches students to design platform teams that actually work.

THEORETICAL FRAME

Skelton and Pais's Team Topologies provides the most operationally useful framework: four team types (stream-aligned, enabling, complicated-subsystem, and platform) and three interaction modes (collaboration, X-as-a-service, facilitating). The key insight: the interaction mode is as important as the team type. A platform team that requires collaboration (negotiated prioritization) rather than providing X-as-a-service (self-serve) has not achieved platform status.

SUGGESTED DISCUSSION ARC

0–10 min

Open with Conway's Law applied to platform teams: 'What does a product look like when it is built by a team that depends on a platform team for every core capability?' Surface: the product's architecture will mirror the interface between the product team and the platform — typically visible as inconsistency (every team used the platform slightly differently) or slowness (every feature required a platform change that took weeks to prioritize).

10–25 min

Lecture: the four platform failure modes. (1) The bottleneck platform: every product team needs something, the platform team has one roadmap, and conflict over prioritization consumes more time than the platform saves. (2) The overbuilt platform: the platform team builds generic infrastructure that no product team actually wants. (3) The underinvested platform: product teams build shadow infrastructure to avoid the platform, defeating its purpose. (4) The governance platform: the platform team becomes a regulatory body that approves or rejects product team decisions. None of these is a platform — they are anti-patterns.

25–55 min

Design exercise: students design a platform team structure for a fictional product company (briefed in advance). The design must specify: (1) what the platform provides, (2) the interaction mode (X-as-a-service, collaboration, or facilitating), (3) how prioritization is decided, and (4) how the platform team's success is measured. Questions 1–3.

55–70 min

Graduate section: Gawer/Cusumano, Henderson/Clark, MacCormack et al., Baldwin/Clark. Questions G1 and G2.

70–75 min

Close: 'Every platform team should be able to state, in one sentence, what it provides as a service, to whom, and how fast. Can the platform teams in your organization pass that test?'

DISCUSSION QUESTIONS — ALL STUDENTS

- 1 The canonical failure mode for platform teams is becoming a bottleneck: every product team needs something from the platform, and the platform team's roadmap can't accommodate all the requests. What organizational mechanisms prevent this failure? Who has decision authority over the platform roadmap?**

Prevention mechanisms: (1) Ruthless prioritization based on platform adoption — only build what multiple product teams need. (2) An internal service-level agreement: the platform commits to a response time and quality standard for platform requests. (3) Explicit escalation path: when a product team believes the platform is blocking them, there is a clear path to resolution that does not require the CPO to adjudicate every week.

- 2 Skelton and Pais describe 'X-as-a-service' as the ideal interaction mode for a mature platform team: product teams consume the platform like an external service, without needing to coordinate or collaborate with the platform team. What does a platform team need to achieve before it can operate in X-as-a-service mode? What must be true about the platform's API, documentation, and reliability?**

Required: a stable, versioned API that product teams can depend on; comprehensive documentation that enables self-serve adoption; an SLA for reliability and latency; a clear deprecation policy; and a feedback mechanism that routes product team issues to the platform team without requiring a meeting. Most platforms claim X-as-a-service but require collaboration in practice.

- 3 A product organization has built a design system as a shared infrastructure for all its product teams. The design system team has 4 engineers and 2 designers. Product team A has adopted the design system for 90% of their UI. Product team B has adopted 10% and is actively building shadow components. What does Product Team B's behavior reveal, and how should the design system team respond?**

Product Team B's behavior reveals: either the design system doesn't meet their needs (wrong abstraction, missing components, or too slow to evolve) or the adoption cost is too high relative to the benefit. The design system team should first diagnose which: interview Product Team B's designers, map the gap between what the design system provides and what Team B is building in shadow, and decide whether to close that gap or explicitly scope the design system's coverage to exclude Team B's use case.

GRADUATE / MASTER'S QUESTIONS

- G1 Henderson and Clark's 'architectural innovation' describes innovations that change the linkages between components while leaving the components themselves unchanged. Apply this concept to the boundary between platform teams and product teams. When a platform team changes its API (the linkage), what are the consequences for product teams even if the underlying capability is unchanged?**

- G2** MacCormack et al.'s empirical study of software systems finds strong evidence for Conway's Law: systems with more tightly coupled architectures come from organizations with more integrated team structures. If you wanted to build a more modular product architecture, what org design change would you make first? Why must the org change precede the architecture change?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students often design platform teams that are simply 'shared services' — rebranded functional teams (shared QA, shared data engineering) rather than genuine platform teams. Shared services are organized around function; platform teams are organized around the product surface they enable. The distinction determines everything about how they are prioritized and governed.

The 'X-as-a-service' ideal is sometimes read as 'platform teams should require no relationship with product teams.' Strong platform teams maintain deep relationships — they understand their customers (product teams) and evolve the platform proactively. X-as-a-service describes the interaction mode for consuming the platform; it doesn't describe a relationship of indifference.

Graduate students may over-apply the Henderson/Clark architectural framing and treat all platform decisions as architectural decisions. The org design implication is simpler: if changing the platform requires extensive coordination with product teams, it is not yet a true platform — it is a shared codebase managed by committee.

CONNECTIONS TO OTHER SESSIONS

The platform team session applies Conway's Law (Week 1) most directly. It connects to Week 10 (enterprise org design — at large scale, entire business units can function as platforms for other business units) and Week 15 (final presentations — many org design proposals include a platform team component that must be carefully specified). The Skelton/Pais framework (four team types, three interaction modes) is a tool students should apply throughout the course.

FACILITATOR NOTES

The design exercise requires a pre-prepared company brief. Keep it simple: a B2C fintech company, 150 engineers, 30 PMs, three product lines (payments, savings, credit), a data platform that all three need, and a design system that two of three have adopted. The brief should be just detailed enough to make the design exercise concrete without being so complex that it dominates the session.

Week 7 · The Agentic Era

FORMAT

Seminar + futures exercise

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

AI agents — software that can autonomously execute multi-step tasks, make decisions, and interact with other systems — are changing the calculus of product team design in ways that the existing org design literature does not address. When an AI agent can autonomously run A/B tests, draft product specs, triage support tickets, or synthesize user research, what does the PM's job become? How does the Rule of Seven apply when some of your team's 'members' are software? This session does not pretend to have definitive answers — it equips students to ask the right questions.

THEORETICAL FRAME

Hoffman's 'Rule of Seven' (agentic era section) is the primary text — this is the chapter that has no precedent in the org design literature, because the situation it describes is new. Cagan's 'Transformed' provides the adjacent software-PM-operating-model context. Brynjolfsson and McAfee provide the economic baseline. The session should feel genuinely exploratory: the instructor does not know the answers to some of these questions, and should say so.

SUGGESTED DISCUSSION ARC

0–10 min

Open with a provocation: 'If an AI agent can do the work of a mid-level PM — user research synthesis, competitor analysis, spec writing, ticket triage, roadmap documentation — what is the mid-level PM's job?' Let students respond freely. Do not resolve the tension. Surface: the job changes, but the degree and direction of change is genuinely uncertain.

10–25 min

Lecture: three hypotheses about what AI agents do to org design. (1) The compression hypothesis: AI agents do the routine work, enabling each PM to own more scope — spans of control widen, teams shrink, the Rule of Seven applies to fewer people. (2) The elevation hypothesis: AI agents do the execution work, freeing PMs to do the strategic and relational work that agents cannot — but the nature of PM skill shifts toward judgment, stakeholder management, and product taste. (3) The restructuring hypothesis: the PM role disaggregates into a smaller number of higher-skill roles and a larger number of agent-oversight roles. All three hypotheses may be simultaneously true for different types of PM work.

25–55 min

Seminar discussion: Questions 1–3. The Rule of Seven in an agentic organization is the most analytically novel question — take the most time here.

55–70 min

Graduate section: Autor, Acemoglu/Restrepo, Simon, Russell. Questions G1 and G2.

70–75 min

Futures exercise: students write a 3-sentence org design recommendation for a product team that will deploy 5 AI agents in the next 12 months. Share and discuss. There is no

right answer — the exercise is diagnostic of how students are thinking.

DISCUSSION QUESTIONS — ALL STUDENTS

- 1 The Rule of Seven is derived from the cognitive limits of human managers managing human team members. If some of a PM's 'direct reports' are AI agents, does the rule still apply? What changes and what stays the same?**

What stays the same: the manager still has finite attention and cognitive capacity. What changes: agents don't require coaching, mentorship, or emotional support — they require different kinds of oversight (output review, goal calibration, failure mode monitoring). The span of control for mixed human-agent teams is an open empirical question. Tentative hypothesis: managers can oversee more agents than humans, but the nature of the oversight work is different.
- 2 If AI agents can do the routine work of PM — research synthesis, spec writing, ticket triage — what does the 'residual PM skill' become? What human capabilities remain irreplaceable at the PM level in a world of capable AI agents?**

Candidates for irreplaceable PM capability: judgment about what to build and why (product taste), relationships with customers, partners, and internal stakeholders, navigation of organizational politics, ethical judgment about product decisions, and the accountability that only a named human can bear. Students should be able to defend their list against the counter-argument: 'but AI will get there.'
- 3 A product organization deploys AI agents that can autonomously run user research, write specs, and generate A/B test hypotheses. How does this change the org design? Specifically: does the team need more PMs, fewer, or different ones? Does the management layer change?**

There is no settled answer. The most defensible position: fewer junior PMs (agents handle the execution work), more senior PMs with broader scope, and a new role (agent overseer / AI product manager) that monitors agent performance and calibrates agent goals. The management layer probably flattens — if agents handle the routine, managers spend more time on the work that requires human judgment.

GRADUATE / MASTER'S QUESTIONS

- G1 Simon's 'Architecture of Complexity' argues that complex systems are hierarchically organized because hierarchy reduces the computational cost of managing complexity. AI agents are being introduced as a new layer in this hierarchy. How do Simon's arguments about hierarchical organization apply to a product team that includes both humans and AI agents? Does the optimal hierarchy change?**
- G2 Acemoglu and Restrepo's empirical research finds that automation that substitutes for labor in existing tasks reduces employment and wages, while automation that creates new tasks increases them. Apply this framework to AI agents in product teams. Are AI agents primarily substituting for existing PM tasks, or creating new PM tasks? What does the answer imply for how product organizations should structure their teams?**

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students sometimes treat the agentic era session as a threat assessment: 'will AI replace PMs?' That is the wrong question. The right question is: how does the PM role change when agents handle routine work, and how should org design respond? The course is not a career advice session — it is an org design session.

The 'AI won't replace strategic work' argument is sometimes stated as though it is settled. It is not. The pace of AI capability development means that work that is 'strategic today' may be routine in 3–5 years. Students should be skeptical of certainty in either direction.

Graduate students may want to apply labor economics frameworks (Autor, Acemoglu/Restrepo) as definitive predictors. Those frameworks are calibrated on historical automation data; AI agents may not follow the same substitution/complementarity patterns. Encourage epistemic humility.

CONNECTIONS TO OTHER SESSIONS

The agentic era session connects back to Week 2 (Rule of Seven — the rule's applicability in mixed human-agent teams) and forward to Week 10 (enterprise org design — large enterprises will face this question at scale before small companies do). It also connects to the final presentation: students whose final proposals ignore the agentic question may be missing a significant structural trend that their proposed org will need to accommodate.

FACILITATOR NOTES

This is the session where the instructor should be most honest about uncertainty. Saying 'I don't know the answer to that' is pedagogically valuable here in a way that would undermine other sessions. The futures exercise at the end should be treated as genuine exploration, not as a test with a right answer. Collect the 3-sentence recommendations and post them — the variety across students will be instructive.

Week 8 · Midterm Case

FORMAT

Presentations + structured critique

TIMING

90 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

The midterm case presentation is the course's first major applied exercise. Students present an organizational diagnosis of a real or simulated company — applying the Rule of Seven, Conway's Law, Greiner's growth model, the mission test, and the platform/product distinction — followed by a redesign recommendation with explicit trade-off analysis. The session is structured around the diagnostic framework from the @arcktip text.

THEORETICAL FRAME

The Amazon two-pizza team model is the session's primary case, but students may choose to analyze their own organization if they have sufficient access to the organizational details. The diagnostic framework has three components: (1) What is the current organizational state? (2) What specific pathologies does that state produce? (3) What is the minimum viable redesign — the smallest structural change that would address the most damaging pathology?

SUGGESTED DISCUSSION ARC

0–10 min

Set expectations for the critique session: the goal of feedback is not to evaluate whether the redesign is 'right' — there is rarely one right answer. The goal is to identify: (1) which course frameworks are being applied rigorously, (2) where assumptions are unexamined, and (3) whether the implementation plan is realistic.

10–70 min

Presentations: each student presents for 8 minutes + 7 minutes structured critique. Critique protocol: (1) What pathologies are diagnosed, and are they traced to structural causes? (2) Which course frameworks are applied, and are they applied correctly? (3) What assumptions are baked into the redesign? (4) What could go wrong in implementation?

70–85 min

Plenary synthesis: 'Across these diagnoses, what is the most common pathology? What is the most common error in redesign recommendations?'

85–90 min

Close: 'The three most important things a good org diagnosis must show are: the specific mechanism by which the current structure produces the observed pathology, the logic connecting the redesign to the reduction of that pathology, and the trade-offs the redesign creates.'

DISCUSSION QUESTIONS — ALL STUDENTS

- 1 Amazon's two-pizza team rule (teams small enough to be fed by two pizzas) is a span of control heuristic. Apply the Rule of Seven to Amazon's team structure: what is the implied maximum team size? What are the organizational consequences of that maximum at Amazon's scale?**

Two pizzas ≈ 6–8 people. At Amazon's scale (100,000+ employees in tech), this implies many thousands of teams and many layers of coordination. The consequence: Amazon needs extraordinarily strong coordination mechanisms at the layer above the team — the narrative memo culture, the six-page memo, the PR/FAQ — to compensate for the coordination costs of thousands of autonomous small teams.

- 2 A diagnosis of organizational dysfunction should trace a specific pathology to a specific structural cause. Practice: 'This organization cannot make decisions quickly because _____. ' Complete that sentence in a way that identifies a specific structural cause (not a cultural cause).**

Examples of structural (not cultural) causes: decision rights are unclear (not: 'the culture doesn't empower people'); the approval chain is too long (not: 'leadership doesn't trust the team'); the team that must make the decision doesn't have access to the information it needs to make it (not: 'silos'). Culture diagnoses are often proxies for structural diagnoses.

- 3 What makes the 'minimum viable redesign' the right framing for an org design recommendation? Why not recommend the optimal redesign?**

The minimum viable redesign is right because: (1) every org change has switching costs — teams must re-learn who they work with, managers must re-establish relationships, information must be re-routed; (2) the second-order effects of large changes are harder to predict than those of small changes; (3) a smaller change is easier to execute, monitor, and reverse if it doesn't work. The optimal redesign on paper is often the wrong redesign in practice.

GRADUATE / MASTER'S QUESTIONS

- G1** Nadler and Tushman's congruence model predicts that organizational performance is a function of the fit between four elements: work, people, structure, and culture. Apply the congruence model to the Amazon two-pizza team case. Which elements are in high congruence? Which are in tension?
- G2** Weisbord's six-box model asks diagnosticians to assess purposes, structure, relationships, rewards, leadership, and helpful mechanisms. Which of these six boxes is most commonly identified in org diagnoses, and which is most commonly overlooked? Why?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students often diagnose cultural problems when the underlying cause is structural. 'The team doesn't communicate well' is a cultural diagnosis; 'the team's interface has no defined owner and the two sub-teams have conflicting incentives for its design' is a structural diagnosis. The distinction matters because structural problems have structural fixes; cultural problems diagnosed as structural are fixed incorrectly.

The midterm redesign is sometimes proposed as a full organizational restructuring. The minimum viable redesign principle pushes back: what is the smallest change that addresses the most important pathology? A full restructuring proposed for a mid-size company in a 10-minute presentation has not been adequately analyzed.

Graduate students may over-apply the congruence or six-box models as checklists rather than diagnostic lenses. The goal is not to fill in all the boxes — it is to identify which box contains the primary pathology and trace the mechanism through which that pathology propagates.

CONNECTIONS TO OTHER SESSIONS

The midterm case is the diagnostic practice ground for the final presentation (Week 15). Students who struggle with the diagnostic framework here have five weeks to develop it further. Connect the midterm feedback explicitly to the final: 'The gap identified in your midterm critique is what you must close in your final presentation.'

FACILITATOR NOTES

Critique pairs work best when students have studied different companies — the cross-company comparison often surfaces insights that a single-case analysis misses. Assign pairs before class and ask each critic to prepare one specific question about the structural mechanism connecting pathology to cause. That preparation makes the critique session much more productive than open-ended feedback.

Week 9 · Cross-functional Collaboration**FORMAT**

Seminar + role simulation

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

The triad of product manager, designer, and engineer is the fundamental unit of modern product development. How that triad is structured — who has final decision authority, how disagreements are resolved, what the career path looks like for each role — determines the quality and velocity of everything the product organization ships. This session addresses the structural and psychological preconditions for high-performing cross-functional teams, using Edmondson's psychological safety research as the empirical foundation.

THEORETICAL FRAME

Psychological safety is not 'niceness' — it is the belief that you can speak up with ideas, questions, concerns, or mistakes without fear of repercussions. Edmondson's research shows that psychologically unsafe teams systematically under-report problems, which means managers make decisions based on filtered, optimistic information. In product teams, that pattern produces products that don't solve real user problems (the user research was filtered), don't get shipped (the technical risks were underdisclosed), and don't generate real business impact (the business case was sanitized).

SUGGESTED DISCUSSION ARC**0–10 min**

Open with Edmondson's hospital nursing study: better-performing nursing teams reported more errors, not fewer, because they felt safe to report errors rather than hiding them. The worse-performing teams had fewer errors in the data — not because they made fewer, but because nurses were afraid to report them. Ask: 'Where in a product organization does the same dynamic occur?' Surface: user research that confirms hypotheses rather than challenging them; technical risk assessments that understate problems; business case projections that hit the target because the analyst knew what answer was wanted.

10–25 min

Lecture: the four structural conditions for cross-functional psychological safety. (1) Decision rights are clear — everyone knows who makes the final call on what, so disagreement is about substance rather than authority. (2) Mistakes have named owners — when something goes wrong, there is a clear accountability path that does not require collective blame. (3) Information is shared symmetrically — all three roles (PM, designer, engineer) have equal access to user research, business context, and technical constraints. (4) Escalation is designed to be safe — there is a clear, low-cost path for raising concerns above the team level without retaliation.

25–55 min

Role simulation: triads of students play PM, designer, and engineer in a fictional product decision. The situation: the PM wants to ship a feature that the designer believes has significant usability problems and the engineer believes has significant technical risk. The simulation runs for 15 minutes. Debrief using Questions 1–3.

55–70 min

Graduate section: Ancona/Caldwell, Eisenhardt/Tabrizi, Katzenbach/Smith, Argyris.

Questions G1 and G2.

70–75 min

Close: 'The highest-leverage action a product leader can take to improve cross-functional collaboration is not a process change — it is demonstrating their own willingness to be wrong in public. What does that look like in practice?'

DISCUSSION QUESTIONS — ALL STUDENTS

- Edmondson's psychological safety research shows that team leaders who model fallibility — who admit when they don't know, who respond to bad news with curiosity rather than punishment — produce more psychologically safe teams. What structural mechanisms reinforce that leadership behavior? What structures undermine it?**

Reinforcing structures: blameless post-mortems (separate learning from accountability), open user research sessions (everyone sees the same data), explicit 'pre-mortem' practices (imagine the project failed — what went wrong?). Undermining structures: management-by-exception (only escalated problems are visible), stack ranking (which makes colleagues competitors), and status hierarchies that penalize engineers or designers for speaking up to PMs.

- In a cross-functional product team, who should have final decision authority over the product roadmap? Over the product design? Over the technical architecture? What happens when those three decision domains conflict?**

Framework: the PM has final authority over what to build and why (roadmap and prioritization); the designer has final authority over how the user experience works; the engineer has final authority over how it is built (technical architecture). Conflict: when the PM's 'what' requires a design approach the designer believes is wrong, or a technical approach the engineer believes creates unacceptable debt, there is no clean authority boundary. Resolution: the team's shared outcome metric (the mission) is the tiebreaker. If the conflict cannot be resolved at the team level, it escalates to the manager — but that escalation path must be cheap, or teams will avoid using it.

- A product team has a 'dominant engineer' — an engineer whose technical opinions consistently override the PM and designer in team decisions, even on product and design questions. What structural changes would address this imbalance? What should the team's manager do?**

Structural changes: explicit decision rights documentation, separate retrospectives for product, design, and technical decisions (to surface patterns), and potentially separating the technical lead role from the team's informal authority figure. Manager actions: one-on-ones that surface the pattern with the dominant engineer directly; a team conversation about decision rights that names the pattern without naming the person; and reinforcing PM and designer authority by asking explicitly for their recommendation in team meetings before the engineer speaks.

GRADUATE / MASTER'S QUESTIONS

- G1 Eisenhardt and Tabrizi find that cross-functional product teams that use an 'experiential' process (building, testing, iterating) outperform teams that use a 'compression' process (planning, coordinating, executing) in uncertain environments. What organizational conditions enable experiential process? What conditions force teams into compression mode even when experiential**

process would be better?

- G2** | Argyris's double-loop learning describes organizations that can question and change their underlying assumptions, not just their actions. What organizational structures enable double-loop learning in a cross-functional product team? What structures prevent it? Is psychological safety sufficient, or is it necessary but not sufficient?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Psychological safety is frequently confused with comfort. Edmondson is explicit: psychological safety is not the absence of conflict, challenge, or accountability — it is the presence of the belief that speaking up is safe. A psychologically safe team can have intense debates; a comfortable team avoids them.

The 'who owns the decision' question is sometimes treated as purely process. It is also political: decision rights determine who has power in the organization. Giving designers final authority over UX is a political act that shifts power from PMs and engineers. Name that explicitly.

Graduate students may treat Argyris's double-loop learning as a diagnosis tool (does this org do it?) without connecting it to specific structural interventions. The question is not only 'does this org do double-loop learning?' but 'what would need to change structurally for it to do so?'

CONNECTIONS TO OTHER SESSIONS

Cross-functional collaboration is the session that most directly addresses the interpersonal and psychological dimensions of org design — dimensions that the Rule of Seven and Conway's Law do not capture. It connects to Week 13 (mentorship — psychological safety is a leadership-created condition) and Week 14 (leading the redesign — org changes are the moments of maximum psychological unsafety, and must be handled accordingly).

FACILITATOR NOTES

The role simulation requires pre-prepared character cards for the PM, designer, and engineer. Each card should specify: (1) what the character knows that others don't, (2) what the character is afraid to say, and (3) what the character wants out of the decision. The best simulations are those where the character's fear maps to a realistic structural incentive — the designer is afraid to speak up because the PM controls their performance review; the engineer is afraid because the last time they flagged technical risk, the PM shipped anyway. Those structural asymmetries are the real subject of the debrief.

Week 10 · Enterprise Org Design

FORMAT

Seminar + guest

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

Managing hundreds of PMs across multiple product lines, business units, and geographies is a qualitatively different problem from managing twenty. The enterprise CPO's organizational challenges — portfolio structure, P&L ownership, career path architecture, cross-business-unit coordination, and the relationship between product and commercial — have no equivalent at startup scale. This session uses the specific organizational challenges of enterprise product leadership, grounded in the Rule of Seven applied at massive scale.

THEORETICAL FRAME

Lafley and Martin's 'Playing to Win' provides the portfolio strategy context: at enterprise scale, org design and product strategy are inseparable — you design the org to execute the strategy, and the strategy is constrained by what the org can execute. Galbraith's work on matrix structures is the technical reference; Goold and Campbell's nine design tests provide the diagnostic toolkit.

SUGGESTED DISCUSSION ARC

0–10 min

Open with a concrete challenge: 'You are the CPO of a 500-person product organization with 80 PMs across 6 product lines. Three of the product lines serve the same enterprise customer. Two of the product lines have overlapping technical infrastructure. One product line is shrinking, one is growing at 40% per year. Design me an organizational structure for this portfolio.' Surface: there is no right answer — but the session's frameworks will help students reason toward a defensible one.

10–25 min

Lecture: three enterprise org design models. (1) Business unit structure: each product line has its own P&L, PM team, and engineering team — maximum autonomy, maximum duplication. (2) Functional structure: PMs, engineers, and designers are organized by function and assigned to products — maximum efficiency, minimum team ownership. (3) Matrix structure: PMs and engineers are assigned to product lines but also belong to functional communities — maximum flexibility, maximum coordination cost.

25–55 min

Case analysis or guest practitioner conversation: a senior product leader from an enterprise company on the specific org design decisions they made and their consequences. If no guest: case analysis of a specific enterprise product org (Netflix, Google, Microsoft, or Meta). Questions 1–3.

55–70 min

Graduate section: Galbraith, Goold/Campbell, Prahalad/Doz, Bower. Questions G1 and G2.

70–75 min

Close: Goold and Campbell's nine design tests applied briefly to the morning's case. 'Which of the nine tests does this organization's structure fail? That failure is the most

important org design problem to solve.'

DISCUSSION QUESTIONS — ALL STUDENTS

- 1 **At enterprise scale, the CPO faces a portfolio management problem: some product lines should grow, some should be maintained, some should be harvested, and some should be discontinued. How does the org design reflect that portfolio strategy? What happens to the team assigned to a product line that the company has decided to harvest?**

Portfolio structure maps to team investment: growth products get dedicated, empowered teams with deep PM and engineering talent; maintenance products get smaller, more operationally focused teams; harvest products may be handed to a 'managing the decline' team or folded into a platform. The org design must make these differences explicit, or the company will under-resource growth products while over-investing in declining ones.

- 2 **The matrix organization is the most common structure for large product organizations and also the most dysfunction-generating. What specific dysfunctions does the matrix produce? Under what conditions is a matrix worth the coordination cost?**

Matrix dysfunctions: unclear accountability (the functional manager and the product manager both have authority — who wins?), priority conflicts (the functional manager wants the engineer to work on platform investment; the product manager wants the feature), and career path confusion (does the engineer report to the engineering VP or the PM?). Worth the coordination cost when: the organization needs both deep functional excellence and strong product integration simultaneously — the condition that typically exists in enterprise companies with complex products across many verticals.

- 3 **A product organization of 500 PMs spans 6 countries. The CPO wants to create a unified PM career path and compensation structure globally. What organizational changes does that require, and what resistance will the CPO face?**

Required changes: centralized PM job architecture (leveling criteria, titles, competencies), centralized compensation bands (which will reveal pay inequities across geographies), and a centralized performance review process. Resistance: country-level leaders who have customized career paths for local talent markets; finance teams that have structured local compensation separately; and PMs who will lose in the harmonization. The CPO must have executive sponsor support before beginning.

GRADUATE / MASTER'S QUESTIONS

- G1 Galbraith's matrix structure analysis suggests that matrices fail most often when the axes of the matrix have incompatible objectives — the geography axis optimizes for local responsiveness while the product axis optimizes for global consistency. How do you design a matrix where the axes have compatible rather than competing objectives? What does that require in terms of incentive alignment?
- G2 Bower's 'Managing the Resource Allocation Process' argues that organizational structure determines which products get resources more powerfully than any explicit strategy process. At enterprise scale, how does the CPO ensure that the resource allocation process reflects the portfolio strategy rather than the political power of each product leader?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students often treat the matrix as a structure to be designed rather than a trade-off to be managed. No matrix eliminates the underlying tension between functional depth and product focus — it just changes where the tension is felt. The question is not 'how do we design a matrix that works?' but 'how do we manage the inevitable tensions a matrix creates?'

The enterprise-startup contrast is sometimes framed as: enterprises are bureaucratic and slow, startups are agile and fast. The more accurate framing: enterprises have coordination challenges that startups don't have, and must design more explicitly for coordination. Startup structures fail predictably at enterprise scale; the question is how to evolve the structure without losing the culture.

Graduate students may apply Galbraith's Star Model as a design tool (fill in all five elements) rather than as a diagnostic tool (identify which element is the primary source of misalignment). The model's value is diagnostic, not prescriptive.

CONNECTIONS TO OTHER SESSIONS

Enterprise org design is the session that applies every previous framework at the largest scale. It connects directly to the final presentation (Week 15) for students who are analyzing enterprise companies. The Week 11 guest session is most valuable when the guest is an enterprise product leader who can speak to the specific challenges discussed here.

FACILITATOR NOTES

If a guest practitioner is available, this is the session for them — the enterprise CPO challenge is difficult to make concrete without firsthand experience. The guest should be prepared with a specific org design decision they made: what the structure was before, what they changed, why, and what happened. The most instructive guests are honest about what didn't work as planned.

Week 11 · **Executive Guest Session**

FORMAT	TIMING	LEVEL
Structured interview + Q&A	75 min	MBA + EMBA

SESSION OVERVIEW

The executive guest session is the course's primary practitioner window. A senior product leader — ideally a CPO, VP of Product, or product founder with experience at multiple company stages — shares the organizational decisions that most shaped (or constrained) their product roadmap. The session is structured as a facilitated interview rather than a presentation, ensuring that specific decisions and their organizational consequences are examined rather than curated success narratives.

THEORETICAL FRAME

The session's primary texts (Edmondson Ch. 6–8, McChrystal Ch. 6–8) are assigned to graduate students as preparation for asking better questions — questions that probe the organizational conditions the guest created or failed to create. The instructor prepares 4–6 questions in advance; student questions are integrated throughout rather than reserved for a Q&A block.

SUGGESTED DISCUSSION ARC

0–10 min	Brief framing by the instructor: 'Our guest today has made organizational decisions that most of us have not had to make yet. We are going to use the course's frameworks to understand those decisions — not to evaluate whether they were right, but to understand the org design logic that produced them.'
10–50 min	Structured interview. Suggested questions: (1) What was the organizational structure of your product team when you joined, and what was the first structural decision you made? Why? (2) Tell me about a team that was performing well — what was it about the team's structure or conditions that made it work? (3) Tell me about an organizational decision that didn't go the way you expected — what did you learn about what org design can and cannot do? (4) How has the agentic era (AI tools in product work) changed how you think about team structure and size?
50–65 min	Student questions: student-prepared questions, integrated into the conversation. Instructor moderates to ensure depth over breadth — one good follow-up question is worth more than three new questions.
65–75 min	Close: instructor asks the guest: 'What is the single most important piece of organizational advice you would give to a new CPO at a 100-person company?'

DISCUSSION QUESTIONS — ALL STUDENTS

1 [For the guest] You used the phrase [quote from interview]. What organizational structure made that possible, and what would have prevented it?

This question is prepared dynamically during the interview — identify the most interesting phrase the guest uses and trace it back to organizational conditions. This is the most valuable question type in a practitioner interview.

2 [For students, post-interview] Which moment in the conversation most clearly illustrates the Rule of Seven? Where did the guest's organization violate the rule, and what did it cost?

Debrief question — apply the course frameworks to the interview content. Students should be prepared to do this in real time.

3 [For students, post-interview] What did the guest say that was most different from what the readings predict? Where did the practitioner account complicate or contradict the theory?

The gap between theory and practice is the most instructive data point from any practitioner session. Protect time to surface it.

GRADUATE / MASTER'S QUESTIONS

G1 Apply Edmondson's psychological safety framework to the guest's account. Where did the guest create conditions of psychological safety? Where did they describe its absence? What structural mechanisms were responsible for each?

G2 Apply McChrystal's shared consciousness / empowered execution framework to the guest's organizational account. How did information flow across the guest's org? Who was empowered to act? Where did the organization fail at either shared consciousness or empowered execution?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students sometimes treat the guest session as a source of prescriptions ('the guest said you should always do X'). The guest's experience is evidence, not instruction. The same org design decision can produce different results in different contexts — the skill is understanding why, not copying what.

Executive guests sometimes present curated success narratives. The instructor's job is to probe for the decisions that didn't go as planned, the trade-offs that were made, and the things the guest would do differently. Those are the most instructive moments.

CONNECTIONS TO OTHER SESSIONS

The guest session is the course's direct practitioner link — it connects every theoretical session to lived organizational experience. It is especially valuable if the guest has experience relevant to the final presentation companies students are analyzing. Ask the guest at the close: 'Which of the companies our students are analyzing faces an org design challenge you have personally encountered?'

FACILITATOR NOTES

Guest preparation is critical. Send the guest: (1) a one-page course overview, (2) the four prepared interview questions, (3) a request to come prepared with one specific organizational decision — not a success story, but a decision that had a specific consequence (expected or unexpected) — and (4) an assurance that confidential details can be shared with the room.

Week 12 · Mergers & Reorgs

FORMAT

Seminar + case

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

Organizational change is where org design theory meets organizational reality — and where it most commonly fails. The merger and reorg context is the harshest test of an org design because it must be implemented against active resistance, within a time constraint, and in an environment where the organizational culture is under simultaneous stress. This session addresses both the strategic org design questions (what should the combined organization look like?) and the implementation questions (how do you get there from here?).

THEORETICAL FRAME

The Netflix Culture Deck is the course's most important primary source for this session: it is a document about how Netflix designed its culture to withstand organizational change (high talent density, freedom and responsibility, no brilliant jerks). The Kanter reading provides the historical context for why mergers so rarely achieve their organizational design intentions. Hoffman's @arcktip account of the JUMP/Uber integration provides a first-hand practitioner case.

SUGGESTED DISCUSSION ARC

0–10 min

Open with a question: 'A company you work for is acquired. Your acquirer has a fundamentally different product culture — more process-heavy, more committee-driven, more risk-averse. What happens to your product org in the 12 months after the acquisition? What can the acquired team's leadership do to protect what made it work?' Surface: the culture of the acquirer almost always wins — unless the acquired team has explicit organizational features that make them valuable enough to preserve.

10–25 min

Lecture: the three stages of org integration in M&A. (1) Stabilization (0–90 days): retain key talent, maintain operational continuity, resist the impulse to impose the acquirer's processes. (2) Integration design (90–180 days): determine what is integrated vs. what is preserved, map the new reporting structure, resolve role ambiguities. (3) Steady state (180+ days): new structure becomes operational, culture integration begins. The most common mistake: moving from Stage 1 to Stage 3 without Stage 2.

25–55 min

Netflix Culture Deck analysis: what specific organizational features does the deck describe? Which of those features would survive a merger with a more traditionally managed acquirer? Which would not? Questions 1–3.

55–70 min

Graduate section: Haspeslagh/Jemison, Buono/Bowditch, Schein, Puranam/Srikanth. Questions G1 and G2.

70–75 min

Close: 'The JUMP/Uber integration case (Hoffman @arcktip) illustrates what happens when a highly autonomous, mission-driven product team is absorbed into a large platform

company with different culture and priorities. What org design features of JUMP would you have specified as non-negotiable in the acquisition agreement?

DISCUSSION QUESTIONS — ALL STUDENTS

- 1 **The Netflix Culture Deck describes 'freedom and responsibility' as the core mechanism: Netflix gives employees unusual freedom (no vacation tracking, expense reporting based on judgment) and holds them to unusual accountability (the 'keeper test'). Which elements of that mechanism could survive a merger with a more traditional company? Which would be immediately eroded?**

Survivable elements: high talent density (if the acquirer agrees to maintain hiring standards), context-setting by leadership (depends on whether original leaders are retained). Immediately eroded: process-minimal culture (the acquirer's legal, finance, and HR teams will impose their processes), the keeper test (the acquirer's HR will not permit the same termination philosophy), and the high-trust expense culture (the acquirer's finance controls will override it).

- 2 **Haspeslagh and Jemison distinguish between acquisitions that require 'strategic interdependence' (the acquirer needs to integrate the acquired company's capabilities) and those that require 'organizational autonomy' (the acquirer must preserve the acquired company's culture and ways of working). Apply this framework to a specific acquisition you know. Was the integration strategy aligned with the strategic rationale?**

Classic misalignment: a large traditional company acquires an innovative startup for its culture and speed, then integrates it into the parent's processes, eliminating the very qualities that made the acquisition valuable. The Haspeslagh/Jemison framework predicts this failure: 'capability transfer' acquisitions require organizational autonomy; 'operational restructuring' acquisitions require strategic interdependence. Treating a capability transfer as an operational restructuring is the most common M&A organizational mistake.

- 3 **A product organization is undergoing a reorg — not an M&A, but an internal restructuring that moves 3 teams from one VP to another. The rationale is strategic alignment: the 3 teams' roadmaps are more closely related to the new VP's business than the old VP's. What organizational risks does this reorg create, and what does the company need to do in the 90 days before and after to mitigate them?**

Risks: talent flight (the 3 teams' most senior members may not want to work for the new VP), coordination disruption (the 3 teams have existing working relationships with teams under the old VP that will become cross-VP dependencies), and mission confusion (the 3 teams' existing missions may not align with the new VP's portfolio). Mitigations: announce with clear rationale before the change takes effect; involve the affected PMs in designing the integration; and maintain existing team missions during a transition period before revising them.

GRADUATE / MASTER'S QUESTIONS

- G1 **Schein argues that organizational culture operates at three levels: artifacts (visible structures and processes), espoused values (stated beliefs), and basic assumptions (unconscious beliefs that drive behavior). Which level is most damaged by a merger? Which level is most resilient? What does that imply for how an acquirer should approach cultural integration?**

- G2** Puranam and Srikanth find that acquisitions destroy the acquired company's ability to coordinate tacit knowledge — the informal understandings between colleagues about how to work together — even when they successfully transfer explicit capabilities. What organizational design features could preserve tacit coordination post-acquisition? Is it possible, or is the destruction of tacit coordination an inevitable consequence of integration?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students often assume that culture is the primary risk in M&A. Structure is equally important: if the acquired team's reporting lines, decision rights, and mission boundaries are dissolved immediately, the culture will not survive regardless of the acquirer's intentions.

The Netflix Culture Deck is sometimes read as a universal model. It worked for Netflix in a specific competitive environment (streaming wars, talent war for engineers) at a specific scale (large enough to absorb the costs of high attrition, small enough to maintain selective hiring). Students should be skeptical of its generalizability.

Graduate students may treat the Haspeslagh/Jemison framework as definitive. The framework describes a strategic logic for integration; the political reality of M&A often overrides strategic logic. The acquirer's middle management (who fear the acquired team) can undermine a strategically correct integration plan.

CONNECTIONS TO OTHER SESSIONS

Mergers and reorgs connect directly to Week 14 (leading the redesign — the implementation challenges of a reorg are the same whether it is M&A-driven or internally driven). The Netflix Culture Deck also connects back to Week 3 (mission-driven teams — Netflix's culture operates through mission clarity, not through rules). The Hoffman JUMP/Uber integration case threads through the course as a recurring practitioner example.

FACILITATOR NOTES

If any student has personal experience with a merger or significant reorg, invite them to share it briefly before the case discussion — even 3 minutes of firsthand experience is more instructive than the readings for the specific organizational dynamics. The best facilitator move: after the student shares, ask 'what organizational mechanism produced that outcome?' That question forces the diagnostic frame onto a concrete experience.

Week 13 · Mentorship as Leverage**FORMAT**

Seminar + coaching workshop

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

Senior product leaders often think of their job in terms of decisions made, strategies set, and roadmaps approved. The highest-leverage activity available to a senior leader — investing in the development of their PMs — is systematically undervalued because its returns are invisible, delayed, and attributed to the PMs rather than to the leader who developed them. This session makes the case for mentorship and coaching as an organizational design strategy, grounds it in the Radical Candor framework and the Cagan/Jones coaching model, and gives students structured practice in giving developmental feedback.

THEORETICAL FRAME

The compounding return on PM development is the session's central argument. A PM who improves at product sense by 10% per year grows to be twice as good in seven years. A senior leader who develops five such PMs has created five times the leverage of a senior leader who does the same work themselves. The org design implication: the most important thing a CPO or VP of Product can do is not make great product decisions — it is create the conditions for their PMs to make great product decisions.

SUGGESTED DISCUSSION ARC**0–10 min**

Open with a question: 'Think of the best manager you have ever had. What specifically did they do — not in general, but in specific instances — that developed you?' Collect 5–7 responses. Surface: the most common answers are not 'they gave me autonomy' or 'they trusted me' — they are specific behavioral acts: 'she pushed back on my analysis until I found the real problem'; 'he gave me feedback that no one else had given me and it was exactly right.'

10–25 min

Lecture: three levels of PM development investment. (1) Information transfer: telling the PM what to do or how to do it. This creates dependency, not development. (2) Feedback: giving the PM specific, honest, timely feedback on their work. This builds skill. (3) Coaching: asking questions that help the PM develop their own judgment. This builds judgment, which scales. Most managers operate at Level 1; high-leverage managers operate at Level 3.

25–55 min

Coaching workshop: students practice giving developmental feedback using the Radical Candor framework — care personally, challenge directly. Students bring a real development challenge (a real PM they manage or have managed, or a peer) and receive coaching from the instructor and peers. Questions 1–3 weave through the workshop.

55–70 min

Graduate section: Kegan/Lahey, Ibarra, Dweck, Drucker. Questions G1 and G2.

70–75 min

Close: 'What is one specific developmental investment you will make in a PM you manage in the next 30 days? Name the PM, the skill, and the specific action.' Public commitment — students hear each other's commitments, which increases follow-through.

DISCUSSION QUESTIONS — ALL STUDENTS

- 1** Scott's 'Radical Candor' framework has two axes: caring personally and challenging directly. The two failure modes are 'ruinous empathy' (caring but not challenging) and 'obnoxious aggression' (challenging without caring). Which failure mode is more common in product organizations, and why? What does the more common failure mode cost the organization?

More common in product organizations: ruinous empathy. The typical product manager's social and professional background selects for high emotional intelligence and conflict avoidance. This produces a management culture where underperformance is not named, difficult feedback is not given, and PMs fail at the next level of challenge without having been told they were at risk.

- 2** Cagan and Jones argue that the product leader's most important job is coaching — spending significant time in one-on-ones helping PMs develop product sense, analytical rigor, and stakeholder skill. What gets in the way of that investment at most organizations? Is the barrier a management skill problem or a structural problem?

Both. Management skill: most product leaders were promoted for product skill, not coaching skill, and have not developed a coaching practice. Structural: the product leader's calendar is filled with stakeholder management, executive reviews, and cross-team coordination — the time that should go to one-on-ones is consumed by the system. The structural fix: protect one-on-one time on the calendar as a non-negotiable, the way you protect time for board meetings.

- 3** The compounding return on PM development is highest when the investment is made early in a PM's career. What organizational design features enable early-career PMs to develop quickly? What features slow development?

Development accelerators: early-career PMs given real product ownership (not just feature work), access to users without a PM filter, honest feedback on their product decisions, and visible examples of excellent PM work. Development decelerators: early-career PMs assigned to maintenance products with no user contact, managed by PMs who hoard product decisions, and evaluated on output (did they ship?) rather than outcome (did it work?).

GRADUATE / MASTER'S QUESTIONS

- G1** Kegan and Lahey's 'deliberately developmental organization' (DDO) describes companies that embed personal development into the fabric of work — not as a separate HR program, but as the purpose of every team interaction. What would a product organization need to do to become a DDO? Is it compatible with the pressures of a product company facing competitive urgency?

- G2** Ibarra argues that leadership identity develops through action — that leaders become leaders by acting like leaders before they feel like leaders. What organizational design features create the conditions for early-career PMs to act like leaders before they are in leadership roles? What features

prevent it?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

The Radical Candor framework is frequently misread as 'be more direct.' The framework's primary claim is the opposite: most managers fail by being too indirect (ruinous empathy), and the solution is to combine directness with genuine care. Directness without care (obnoxious aggression) is not Radical Candor — it is a different failure mode.

Mentorship is sometimes equated with friendship or sponsorship. Mentorship in the Cagan/Jones sense is specifically developmental: helping the PM see what they cannot yet see about their own work. It requires skill that is distinct from both friendship and sponsorship.

Graduate students may treat the developmental organization literature as aspirational and utopian. The Kegan/Lahey evidence is real: companies that invest heavily in employee development have measurably lower attrition and higher performance. The question is not whether development is worth it — it is how to build a development practice that survives competitive pressure.

CONNECTIONS TO OTHER SESSIONS

Mentorship connects to every session where team quality is implicated — which is all of them. It is especially important in Week 14 (leading the redesign — the most effective organizational changes are led by leaders who have already developed the trust of the people being reorganized) and Week 15 (final presentations — students should evaluate every proposed org structure against the question: what development opportunities does this structure create?)

FACILITATOR NOTES

The coaching workshop is the most skill-building exercise in the course. Students should practice giving feedback on a real situation — not a hypothetical. The most common workshop failure: students give vague, positive feedback ('I thought your analysis was really strong') rather than specific, developmental feedback ('I noticed you didn't test your demand assumption with actual users — why not, and what would it have changed?'). Model the latter explicitly before students practice.

Week 14 · Leading the Redesign

FORMAT

Seminar + stakeholder mapping workshop

TIMING

75 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

Organizational change is the most politically demanding activity a product leader will engage in. The best org design on paper fails if it cannot be implemented — and implementation requires navigating the interests, fears, and informal power of everyone who is affected by the change. This session addresses the implementation dimension of org design: how to diagnose stakeholder interests, build coalitions, sequence the change to minimize resistance, and sustain the new structure after the initial announcement.

THEORETICAL FRAME

Kotter's eight-stage change model is the most widely taught framework for organizational change, and the course applies it critically: Kotter's model describes what must happen, but not the political dynamics that determine whether it happens. Heifetz and Linsky's adaptive change framework provides the political dimension: some organizational problems are technical (they have known solutions that can be implemented without changing people's values or identities) and some are adaptive (they require people to change their commitments, behaviors, and beliefs). Most org redesigns are adaptive, not technical.

SUGGESTED DISCUSSION ARC

0–10 min

Open with the most common org redesign failure mode: 'The new org chart is announced on a Tuesday morning via all-staff email. By Tuesday afternoon, the informal power structure of the old org has begun reasserting itself.' Ask: 'Why does this happen? What does the organizational announcement ignore?' Surface: the announcement changes the formal structure; it does not change the information flows, relationships, and informal authority that constitute the actual organization.

10–25 min

Lecture: three things a successful org redesign must do that an announcement cannot. (1) Create a felt sense of urgency for the change — people must believe the current state is unsustainable, not just that the leader prefers the new state. (2) Build a coalition of people who benefit from the change and will actively support it. (3) Generate early wins — small, visible evidence that the new structure is producing better outcomes than the old one.

25–55 min

Stakeholder mapping workshop: students map the stakeholders in a specific org redesign (their own, a case, or the final project). For each stakeholder: What do they gain from the change? What do they lose? How much influence do they have? How much resistance do they represent? Questions 1–3 weave through the workshop.

55–70 min

Graduate section: Heifetz/Linsky, Pfeffer (Power), Schein (Process Consultation), Nadler. Questions G1 and G2.

70–75 min

Close: 'The implementation of an org redesign is where most of the learning happens. The structure you announce is not the structure you end up with — because the people adapt it, resist it, and reshape it in ways you didn't predict. What does a good redesign leader do with that information?'

DISCUSSION QUESTIONS — ALL STUDENTS

- 1 Kotter argues that the most common mistake in organizational change is declaring victory too soon — the new structure is announced, the immediate crisis passes, and leadership attention moves elsewhere before the change is embedded. In a product organization, what are the signs that a reorg has been successfully embedded vs. merely announced?**

Signs of embedding: the new teams are producing work that could only come from the new structure (cross-team collaboration that the old structure prevented); informal networks have reorganized to match the formal structure; new team missions have been written and are being used as actual decision guides; and the people who resisted the change have either adapted or left. Signs of merely announced: the old coalitions are still forming around the old org chart; informal escalation paths still follow the old reporting lines; and the new teams' first roadmaps look like the old teams' roadmaps.

- 2 Heifetz and Linsky distinguish between technical challenges (clear problem, known solution, implementation is the challenge) and adaptive challenges (problem definition is contested, solution requires behavior change, and the 'solution' threatens someone's identity or interests). Apply this distinction to an org redesign you know or are planning. Is it technical, adaptive, or both?**

Most org redesigns are adaptive: they require managers to cede control they currently hold, teams to take on accountability they currently avoid, and leaders to be transparent about strategic choices they currently obscure. The mistake: treating an adaptive challenge as a technical one — issuing an org chart without building the coalition, framing the problem, and creating the conditions for people to change.

- 3 A product leader wants to redesign their org to create more empowered product teams. The engineering VP is resistant — the change requires engineering managers to cede product decision authority to PMs, which the engineering VP believes will reduce engineering quality. Map the engineering VP's resistance: what are they protecting, and what would need to be true for them to support the change?**

The engineering VP is protecting: engineering quality standards (legitimate concern — PM-led teams sometimes under-invest in technical quality), their own career trajectory (if PMs have more authority, the engineering VP has less), and their team's culture (engineering identity is partly built on technical authority). What would need to be true: the PM-led model must preserve engineering quality standards (perhaps through explicit technical quality OKRs), the engineering VP's role must gain something in the new structure (perhaps responsibility for engineering excellence across all teams), and the change must be implemented with input from engineering leaders, not imposed on them.

GRADUATE / MASTER'S QUESTIONS

- G1 Heifetz and Linsky argue that leaders who attempt adaptive change are regularly 'attacked from the**

'balcony' — subjected to personal attacks that deflect from the substantive challenge. What organizational design features protect a leader who is attempting an adaptive org change? How do you distinguish between legitimate pushback and defensive resistance?

- G2** Pfeffer argues that organizational change is fundamentally a political process — it requires building coalitions, managing information strategically, and accumulating and spending political capital. How does a product leader build political capital for an org redesign without appearing manipulative? What is the ethical line in organizational politics?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Students often believe that a clear, logical rationale for an org change will produce support from stakeholders. Heifetz and Linsky's insight: people resist adaptive change not because they don't understand the rationale but because they do understand it and see that it threatens something they value. The rational case for the change does not address the adaptive resistance.

The stakeholder mapping exercise is sometimes treated as a political manipulation exercise — 'how do I get these people to do what I want.' The deeper frame: understanding what stakeholders value and how the change affects those values is a prerequisite for designing a change that is actually better for the organization, not just for the leader proposing it.

Graduate students may treat Pfeffer's power framework as cynical or immoral. The response: organizational power is a fact, not a moral position. A leader who ignores power dynamics does not avoid them — they simply operate without understanding them, which produces worse outcomes for everyone, including the people they are trying to help.

CONNECTIONS TO OTHER SESSIONS

Leading the redesign is the culminating practical session before the final presentation. Every org design proposal in Week 15 must include an implementation plan that addresses the political dimension — the stakeholder map, the coalition-building strategy, and the early win sequence. Students who have not engaged with the implementation question in Weeks 13 and 14 will produce analytically sound but organizationally undeliverable final proposals.

FACILITATOR NOTES

The stakeholder mapping workshop requires that students bring a real org change they are planning or have planned. EMBA students almost always have one. MBA students sometimes don't — have a rich fictional case ready for students who need it. The most instructive maps are those where the student initially assigns a stakeholder to 'supporter' and then, through the mapping process, discovers that the stakeholder actually has reasons to resist. That discovery is the point.

Week 15 · Final Presentations

FORMAT

Public presentations + external critics

TIMING

120 min

LEVEL

MBA + EMBA

SESSION OVERVIEW

Final presentations are the course's synthesis event. Teams present a complete organizational design proposal for a real or simulated company: a current-state diagnosis grounded in the course's frameworks, a redesign recommendation with explicit trade-off analysis, and an implementation plan that addresses the political and cultural dimensions of the change. External practitioners — ideally a CPO or VP of Product, an organizational designer or management consultant, and a product founder — serve as critics.

THEORETICAL FRAME

No new theory. The session tests whether students can synthesize the course's full toolkit — Conway's Law, the Rule of Seven, the mission test, Greiner's growth phases, the platform/product distinction, psychological safety, the agentic era implications, and Kotter's change management model — into a coherent, politically realistic organizational design proposal.

SUGGESTED DISCUSSION ARC

0–5 min

Frame: 'A good org design proposal must do three things: diagnose the right problem, recommend a structurally sound solution, and describe a politically realistic implementation. The best proposals are honest about what they don't know and what could go wrong.'

5–95 min

Presentations. For a class of 15–20: teams of 3, each team presenting 15 minutes + 10 minutes panel critique. The panel should be instructed to ask one question each per team — and to spend at least one question on the implementation plan.

95–115 min

Panel synthesis: 'What was the most sophisticated diagnostic you saw today? What was the most common gap? Which implementation plan was most politically realistic?'

115–120 min

Instructor close: return to Week 1's organizing question — 'Organizations are products. They can be well or poorly designed.' Return to the students' own examples from Week 1: 'Now that you have the course's full toolkit, how would you diagnose the dysfunction you described in our first session?'

DISCUSSION QUESTIONS — ALL STUDENTS

1 | What is the single most important structural change in your proposal, and what specific

pathology does it address? Why is this change necessary rather than merely helpful?

The 'necessary rather than helpful' distinction is key: necessary means the current pathology is serious enough that not changing is unsustainable. Helpful means it would be nice but the organization could limp along without it. Only necessary changes get implemented — helpful changes get deferred.

2 Apply Conway's Law to your proposed structure: what products will this structure produce? Are those the products the company needs?

The best answer traces a specific product capability to a specific structural feature. 'This structure places the data team adjacent to the customer-facing PM team, which means data insights will influence product decisions before they are finalized rather than after — which we believe will produce features that are more measurably effective.' That is Conway's Law applied prospectively.

3 Who in the organization will resist this change? What specifically will they lose, and what have you done in your implementation plan to address that loss?

Every org change creates winners and losers. The proposal that pretends otherwise has not thought through the implementation. The best answers name specific roles or departments, describe what they lose (authority, resources, visibility, career path), and describe specific mitigation strategies — not just reassurances.

GRADUATE / MASTER'S QUESTIONS

G1 Apply Heifetz and Linsky's adaptive change framework to your proposal. Is your proposed change primarily technical or primarily adaptive? What evidence supports your assessment, and how does that assessment change your implementation plan?

G2 Looking back at the course: which single framework — Conway's Law, the Rule of Seven, the mission test, psychological safety, Kotter's eight stages, the Greiner model — has most changed how you think about organizations? How does that change show up in your final proposal?

COMMON MISCONCEPTIONS & HOW TO ADDRESS THEM

Teams sometimes present the ideal org structure without the implementation plan. An org design proposal without an implementation plan is an academic exercise, not a professional deliverable. Every proposal must include a sequenced set of steps, a timeline, and an assessment of what could go wrong at each step.

The 'restructuring as solution' fallacy: teams sometimes propose a full restructuring when a smaller intervention — a clarified mission, a new decision-rights framework, or a new coordination mechanism — would address the pathology without the switching costs of a full reorg. Evaluate every proposal against the minimum viable redesign principle from Week 8.

External critics sometimes evaluate proposals on strategic merit ('I'm not sure the strategy is right') rather than on org design quality ('given that strategy, does this structure enable it?'). Redirect to the org

design question: the course is not about strategy, it is about how org design enables or constrains strategy.

CONNECTIONS TO OTHER SESSIONS

The final presentation is the destination of the course's full arc: from 'organizations as products' (Week 1) to 'leading the redesign' (Week 14). Every tool introduced in the course — Conway's Law, the Rule of Seven, the mission test, the growth model, the psychological safety framework, Kotter's eight stages — should appear somewhere in a complete final proposal. The best proposals do not cite all of them by name; they apply the underlying logic without the vocabulary.

FACILITATOR NOTES

External practitioners should be briefed with: (1) the course's central frameworks (Conway's Law, Rule of Seven, mission test), (2) the evaluation criteria for the final proposal (diagnosis rigor, structural soundness, implementation realism), and (3) a request to ask at least one question about the implementation plan in every team critique. The best final sessions feel like a real organizational design review — rigorous, direct, and genuinely useful to the teams presenting.

APPENDIX: QUICK-REFERENCE QUESTION BANK BY THEME

These questions can redirect or deepen discussion in any session:

Urban Systems (any session)

- What does this product do to the urban system? What feedback loops does it create?
- What would Jacobs say about this product? What would Bertaud say?
- What does optimizing for the measurable sacrifice that is not measurable?

Equity & Justice (Weeks 5, 8, 11, 13)

- Who benefits from this product in year one? In year five? Are they the same people?
- Which equity failure mode applies here: design, deployment, or displacement?
- What would Sheller's mobility justice framework require of this product as a condition of legitimacy?

Regulation & Governance (Weeks 2, 9, 14)

- Is this regulation alleviating an externality or redesigning the city? What is the difference in this case?
- What does the regulatory text reveal about what the city actually values?
- What would a collaborative — rather than adversarial — regulatory strategy look like?

Data & Measurement (Weeks 3, 10)

- What does this product measure, and what does optimizing for that metric make worse?
- Who generated this data, who owns it, and who can see it?
- What would Loukissas say about applying this dataset in a different city context?

Business Model & Capital (Weeks 6, 12, 14)

- Who internalizes the profits and who absorbs the externalized costs?
- What happens to this product's mission if it is acquired by a larger platform?
- What is the minimum business model that could sustain this product without extracting value from the city?

Practitioner Perspective (Weeks 5, 13, 15)

- What does Patrick's account of building JUMP see that academic readings cannot? What does it miss?
- What would you do differently from the SoBi/JUMP arc, given what you now know?
- If you were advising a city council on whether to permit this product, what three questions would you ask the product team?

This guide is a living document. Faculty who use it are encouraged to annotate, revise, and return observations to Patrick Hoffman (@arcktip) so that future iterations benefit from classroom experience across institutions.